

Industrial Year Report

Software Developer at ISIS Neutron Source

Samuel Jackson
University Of Aberystwyth
slj11@aber.ac.uk

September 10, 2014

Contents

1	Introduction	3
2	Brief Introduction to Neutron Scattering	3
3	Organisational Environment	4
3.1	Organisation of STFC	4
3.2	Organisation of ISIS Neutron Source	5
3.3	Organisation of the Mantid Project Development Team	5
4	Technical and Application Environments	5
5	Description of Job Role and Work Done	6
5.1	Initial Training and Development Work	7
5.2	Indirect Bayes Interface	7
5.3	VESUVIO Calibration	7
5.4	Density of States Algorithm	8
5.5	Improvements to Algorithm History Recording	8
5.6	Development Reports	9
6	Critical Evaluation of Placement	9
	References	11

1 Introduction

This report details the industrial placement year I undertook as part of the Software Engineering MEng course at Aberystwyth university. The position of the placement was a junior role as part of the Mantid data analysis toolkit development team based at the ISIS facility at Rutherford Appleton Laboratory in Harwell, Oxfordshire, owned by the Science and Technologies Facilities Council (STFC) for the duration of a one year contract of employment. This position was then extended by an additional two months to continue with existing and new project commitments.

ISIS is a world leading neutron and muon scattering facility. The facility operates a 800 MeV pulsed proton synchrotron which acts as a source of neutrons for both neutron and muon spectroscopy [1] using the principle of time-of-flight. Neutron and muon spectroscopy is used to probe the structure and dynamics of materials at truly fundamental level.

The Mantid project aims to provide a single, unified application for the analysis of neutron and muon scattering facilities such as ISIS [2]. The project is primarily developed by two teams of developers, one based at ISIS and one at the Spallation Neutron Source (SNS) at Oakridge laboratory in Tennessee, USA.

2 Brief Introduction to Neutron Scattering

In order to fully understand the operation of ISIS and what the Mantid application is used to analyse, a small amount of background knowledge of the techniques used in neutron scattering experiments are required along with some definitions of key terms. All instruments at the ISIS accelerator operate on the time-of-flight principle. This is where the time for a neutron to travel from the source to the detector of an instrument is accurately known. This value when combined with known parameters of the instrument (i.e. the length of the incident and final flight paths) can be used to determine interesting properties about the sample. The raw value measured by an experiment is the time it takes a neutron to reach a detector in microseconds. Typically this data is stored as a histogram with a count of the number of neutrons detected along the y-axis and the time-of-flight in microseconds on the x-axis.

Broadly speaking, neutron scattering can be split into two categories: elastic and inelastic. elastic scattering is where the final energy of a scattered neutron is equal to the energy of the incident neutron, i.e. there is no transfer of energy to or from the sample. Inelastic scattering (which is the technique generally used by instruments belonging to the MSG) is the more complex case where the energy of the incident neutron and the scattered neutron are not equal, i.e. there is a transfer of energy to or from the sample. From this transfer in energy and from known parameters of the instrument an instrument independent scattering function can be defined which encapsulates the properties of the sample. This function is usually denoted as $S(Q, \omega)$, where Q is the momentum transfer and ω is energy transfer [3, 4].

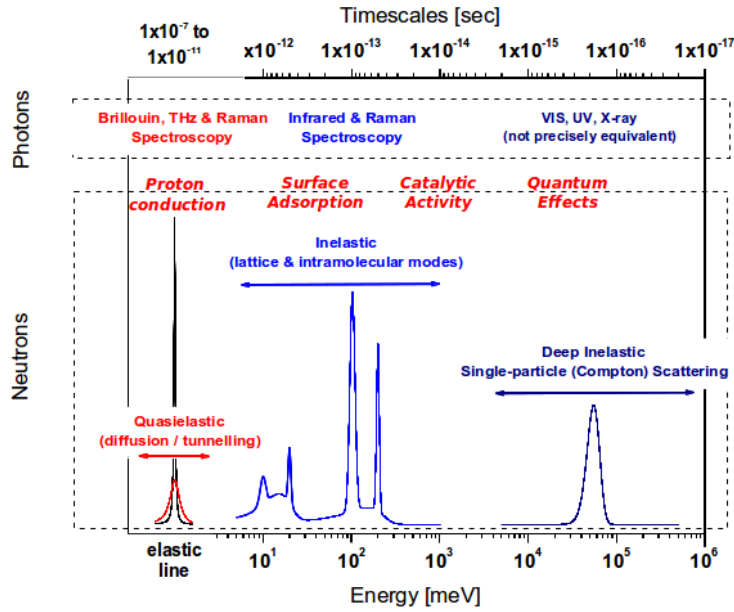


Figure 2.1: Neutron scattering techniques by energy range.

Two important types of inelastic scattering are quasi-elastic neutron scattering (QENS) and deep inelastic neutron scattering (a.k.a neutron Compton scattering). Quasi-elastic neutron scattering is the case where the energy transfer is very close to zero and the scattering is almost elastic. This is typically performed on low energy spectrometers such as IRIS and OSIRIS [5, 6]. Deep inelastic neutron scattering (DINS) is the opposite case which uses extremely high energies ($>1\text{eV}$) and is used to measure the momentum distribution of atoms. This technique is still in its developmental phase, with the VESUVIO spectrometer at ISIS currently being the only instrument in the world capable of doing it [7].

Each of the spectrometers used by the MSG are of what is known as indirect geometry. This is where the final scattered energy is restricted to a particular wavelength (usually with a analyser crystal which will absorb all other wavelengths) and the incident neutron energy is varied. The value of the incident energy can then be found though the laws of the conservation of energy.

My work in with the Mantid team was centred round improving and maintaining the code base of indirect geometry instruments at ISIS. For a more in-depth introduction to neutron scattering the reader is directed to the references, in particular refs. [3, 4, 8] provide the best introductions to those unfamiliar with the subject.

3 Organisational Environment

3.1 Organisation of STFC

The Science and Technologies Facilities Council is a UK government funded body that carries out a wide variety of scientific research across a multitude of disciplines including particle physics, nuclear physics, space science and engineering, medical and biological sciences, and computational science.

While the organisation is funded by the government, it is classified as a non-governmental body which acts as an umbrella organisation for an array of facilities based across the UK. These include (but are not limited to) the central laser facility, diamond light source (which is a publically limited company of which STFC holds an 86% share), ISIS neutron source and RAL space based at Rutherford Appleton Laboratory in Oxfordshire, the Daresbury Laboratory located in Cheshire and the Chilbolton Observatory based in Hampshire.

The organisation has its head office located at Polaris House, Swindon, Wiltshire and is headed by the

Chief Executive John Womersley. The purpose of STFC is to control the general overall management of the facilities under its control, in particular it is responsible for allocating budgetary and staffing allowances and liaising between government departments, particularly the department for business, innovation and skills.

3.2 Organisation of ISIS Neutron Source

ISIS Neutron source is a project that is owned, operated and funded by the STFC. The organisational hierarchy of ISIS is headed by the director Prof. Robert McGreevy and several division heads for individual functional areas within ISIS such as the diffraction, spectroscopy and support, experimental operations, instrumentation, design, and accelerator divisions.

ISIS is also split into a number of research, operations, and experimental support groups. The computing group, of which the Mantid project is a part of, falls into the experimental support group and is technically a part of the scientific computing group which also includes the ICAT data catalogue which manages the data collected from sample runs on the instruments for future analysis.

While this description provides an overview of how the staffing of ISIS can be divided, in practice there tends to be a lot of cross over between sections of the organisation depending on a employees skills and responsibilities. For example, while I was employed as part of the scientific computing group, my line manager and senior manager were both members of the molecular spectroscopy research group whose interests I was responsible for within Mantid.

3.3 Organisation of the Mantid Project Development Team

The Mantid development team in the UK is a subset of the scientific computing group of ISIS. The Mantid group is headed by a single project manager (Nick Draper) who is based at ISIS where the project first started and is responsible for the overall management and direction of the project. The project is split into two teams, one based at ISIS and one at the SNS in Oakridge, Tennessee. Both teams consist of a single lead developer and several senior developers who oversee the major technical developments and help to guide and manage the rest of the development team. Both teams in the UK and the US also have there own manager, but the project manager based at ISIS is in overall control. Within Mantid, the project manager and the majority of the senior developers are actually contractors from Tessella Ltd, based in Abingdon, Oxfordshire. The rest of the development team are directly employed by ISIS, or in the case of the Americans, by the SNS.

While the author was primarily situated within the development team based at ISIS, many developers within the team are generally also attached to a specific scientific group at the facility. In the author's case this was the ISIS Molecular Spectroscopy group (MSG) [9] which specialises in using indirect geometry spectrometers to preform quasi-elastic, inelastic, and deep inelastic (also know as Compton) neutron scattering [3, 4] for condensed matter science and thus much of the work done over the course of the year was carried out in relation to the needs and requirements of the MSG.

4 Technical and Application Environments

The development team based in the UK consisted of a single office located within the main office building of the ISIS facility. This office consisted of approximately 13 workstations for use by the team. The number of machines varied throughout the course of the year depending on the level of staffing available for the project. Each of the machines were reasonably powerful 64-bit Dell workstations with between 8-16 Gb of RAM and 8-16 core intel i7 processors. Typical hard drive space for the machines was between 512 Gb to 1 Tb with the majority still being disk drives, but some of the newer machines had flash based storage.

The operating system that each machine ran was completely left to the preference of the developer, but it was recommended that the developer run one of the operating systems supported by Mantid for obvious reasons. In practice this meant that there was a good variety of developers using different platforms. The author chose to run Ubuntu 12.13 as his operating system of choice for the majority of development work, with a dual partition running Windows 7 which could be swapped to when circumstances required. Other operating systems used by developers in the team included Windows 8, Mac OSX Mountain Lion and Mavericks, Red Hat Enterprise Linux 6, and Fedora 20.

Apart from the workstations, the development also had a collection of Jenkins build servers in order to support a continuous integration and testing workflow in conjunction with the Gitflow workflow [10]. The build servers were jointly located at both ISIS and the SNS. At the start of the placement, the build servers for ISIS and the SNS were completely separate and located at different web addresses. Each individual build was run as a single job on the Jenkins build servers. At the beginning of the current year this was changed so that the servers were located at the same web address and the organisation of the build servers were changed to make use of matrix builds. This is where multiple builds are kicked off at the same time under a single umbrella job. For example the development branch matrix build would build the project and run the unit test on each officially supported OS.

Like the choice of operating system, the development software used by the team was flexible and open to developer preference. The project is built using the CMake build system on all supported platforms. On windows platforms the only supported compiler was Visual Studio 2012 or 2014 and most developers either chose to use the Visual Studio IDE, the Qt Creator IDE, or the Eclipse IDE. On Mac the Intel C++ compiler is used and typical IDEs are XCode, Eclipse, or Qt Creator. On Linux distributions the GNU compiler is the main supported compiler, with either Eclipse or Qt Creator used as the IDE for development. Many Linux developers are also content to just use the make command to build the project from the command line. This approach is often used in conjunction with lightweight editors such as vim or sublime text. For interface development, Qt Designer was used on all platforms for anything more than the most trivial of jobs.

Unit tests are optionally built alongside the project using a separate build target generated by CMake using the CxxTest unit testing framework. System tests are written in python and make use of a collection of custom scripts loosely based on the unittest python module and makes use of the Mantid applications python API. Debugging software used typically makes use of Visual Studio on Windows, XCode on Mac and GDB on Linux distributions.

The Mantid application makes use of data files produced directly from neutron spectrometers. These files are collected on a collection of servers based by ISIS and owned by the scientific computing department and provide both the instrument scientists and visiting scientists with access to their data. The development team also has direct access to all of the data generated from the instruments. This is available through network drives. On Windows operating systems access is provided through the in-built network drive capabilities. On Mac and Linux access is obtained through using Samba software in conjunction with the SMB protocol. Copies of actual instrument data are frequently used as part of test scripts, especially in the case where the data required for the test cannot be easily simulated programmatically.

With regards to project management, the development team makes use of the git version control system and keeps all of the source code openly available to the public via Github. The Trac ticketing program is used to keep a record of the development work that needs doing and for monitoring the progress of development. The Trac setup used also included plugins to automatically capture commits to branches on Github and update the ticket with the appropriate information.

5 Description of Job Role and Work Done

As mentioned in section Introduction, the majority of my year was spent attached to the Molecular Spectroscopy Group (MSG) at ISIS. My role in the development team was to satisfy the computational data analysis requirements of the MSG within Mantid. This included involvement in every part of the

development cycle, from gathering requirements from the users (the instrument scientists in the MSG) through to implementation of requested features, to testing and maintenance/bug fixing.

Before going into discussion about the work I carried out as part of my placement, it is useful to define a couple of core concepts used within the Mantid application: algorithms and workspaces. In Mantid a workspace [11] is an entity that contains a dataset. There are several types of workspace, but the most common is a MatrixWorkspace which is a collection of X, Y, and E spectra either as a histogram or as point data. An algorithm [12] is a routine defined in the Mantid framework that manipulates a workspace in some way (such as loading data into a workspace or rebinning a workspace). A common analogy is that workspaces are the ‘nouns’ of mantid, while algorithms are the ‘verbs’.

5.1 Initial Training and Development Work

My first couple of weeks at ISIS were spent running through some training exercises setup by the development team. These were relatively simple programming problems designed to gauge my existing understanding of C++ and Python. After spending a week completing these introductory activities, I had an informal code review with the lead developer who examined the good and bad points of what I had written before giving me an introduction to the Mantid application and workflow of the team. The following week was mostly spent setting up my workstation for development before starting my first ticket from the Trac ticket system. My first ticket involved writing a couple of algorithms which would take some histogram data and perform smoothing or interpolation of the data using the cubic spline routines in the GSL library, a feature that had been requested by one of the MSG scientists.

5.2 Indirect Bayes Interface

Apart from maintenance and minor feature requests, my first real major piece of development work was to write a GUI for a collection of bayesian fitting routines that had been provided by one of the scientists in the group. These routines are used as part of quasi-elastic neutron scattering analysis to determine the shape of the scattering function present in the dataset by using bayesian methods to compute the likelihood of a given model. These methods were based on the original routines described by Sivia [13].

The purpose of the interface was to provide an easier way for users to set up and run a fit to their data than by using the existing scripts which required many parameters to be specified. This was fiddly and often led to the user inputting incorrect combinations of parameters causing the script to throw an error or even crash.

The solution to this was to add another custom interface under the indirect section of the application (there were already three others in place). This consisted of a collection of C++ classes, created using the Qt framework, one for the parent window of the GUI and one for each of the individual routines on the interface which were implemented as separate tabs on the interface. Each of the routines were refactored from the original code base also required updated system tests, which were implemented at the same time as the new interface.

5.3 VESUVIO Calibration

Another major piece of work I was involved with during my placement was porting some calibration scripts for the VESUVIO instrument [7] originally written as Fortran programs which used the older OpenGenie application [14]. As part of this project, I created a new implementation for Mantid based on the original calibration procedures described in Ref. [15].

This required rigorous testing by the scientist responsible for the instrument and was built slowly over a series of months. The Mantid implementation was designed to be radically different from the original. This version was built using the Mantid concepts of algorithms and workspaces and was intended to

be more modular and more maintainable than the original. A couple of unit test suites used to check the quality of the calibration using several different sample materials were also written as part of the development process.

The end product was a pair of Mantid algorithms, one for fitting the appropriate line shape to each peak in every spectrum of a sample run specifically measured for calibration with the Mantid python API and one which calculated the calibration parameters for the instrument from the fit obtained from the first algorithm.

5.4 Density of States Algorithm

During my time with the MSG, it became apparent that one of the major areas for development that had not been explored was the implementation of support for the simulation of a neutron scattering experiment. This was an area which my supervisor (who has a background in computational simulations) was keen to expand because comparison with simulation is one of the most common techniques for analysing experimental data.

One of her requests was to add support in Mantid for loading the results of simulation produced using the CASTEP program. My supervisor already had a perl script which could do the calculations she required on the command line, but it was infinitely more convenient to have this functionality within Mantid. This involved the creation of a new algorithm which would read the files output from this program, which contain a list of predicted frequencies for a sample. With this information, the density of states could be calculated and a workspace containing the simulated spectrum could be created. Functionality was also added to calculate the infrared and raman spectrum from data in the files. Later in this placement she further requested additional functionality be added to the algorithm to calculate the partial density of states from simulation data.

5.5 Improvements to Algorithm History Recording

Sometimes I was asked or it became necessary for me to work on more general development of the Mantid framework. The largest piece of general development I was asked to do was to improve the algorithm history system. In Mantid, when algorithms get executed on a workspace, a record of the algorithm ran and the properties used to execute the algorithm are stored on the workspace object. In theory this would provide a detailed record of everything that had happened to the workspace from the moment the data entered the application. In practice this feature had only partially been implemented and had slowly become more redundant/broken as development moved on.

The aim was for me to overhaul this and to not only properly provide the features mentioned above, but also to add functionality to capture nested algorithm history, that is, the history of an algorithm operating on a workspace, but also the history of an “child” algorithms run as part of that algorithm. There was also a requirement to make the workspace completely reproducible from its history alone and that the full nested history could be examined using a GUI from within Mantid.

As this was the most core of core changes, and unlike area specific features, a much more rigorous design phase was required. An existing draft of a design document existed for the algorithm history which I presented to the development team (both in the UK and the states) and was roundly rejected for being too vague. The next step was to prototype out some of the more unclear features such as how history should be saved to file which needed to be a compromise space and time efficiency, how to represent nested history at runtime (boost::graph, std::list of children, std::set of children?), and how nested history records should be stored on disk.

After prototyping the system, I created a new design document with the findings from the prototypes and my suggested proposals for change [16]. This document was reviewed by the senior members of the development team and with a few minor changes was approved. I then spent the next month or so implementing the changes to support the requested features. Along the way we encountered several issues that were not covered by the design document, mostly linked to the fact that workspaces can be created

without a string name and that providing a temporary name to them on the fly can cause many issues with reproducing an executable script.

Despite these setbacks, I managed to get the majority of the core changes into Mantid in time for the next release. However there were still several parts of the system left to the next release cycle, such as compressing the size of large algorithm properties (such as when arrays are passed as parameters to an algorithm).

5.6 Development Reports

One final piece of work I was asked to complete during my time at ISIS was to produce two major reports. The first was to provide a snapshot of the current progress of development of the indirect framework. This report gives an overview of the entire framework as of release 2.6 including a description of the theory behind the routines and detailed description of how each one works. Suggested ideas for future development are also listed. The second was a basic user manual for the routines still under development for the VESUVIO spectrometer but provides a description of how to analyse data from the instrument with the implementation as it stands. Both of these reports were finished before I left the company and are due to be published in the near future. In the meantime, a copy of both reports are available on request.

6 Critical Evaluation of Placement

My overall experience on the placement was a positive one. The placement at ISIS provided me with a level of responsibility far exceeding that which I expected. For the most part, I felt that I was a valued member of the development team and that I was treated as an equal rather than as an intern. I feel that my practical (and in some respects theoretical) understanding of my subject area has increased tenfold since, particularly over the first few months which I found to be fairly steep but enjoyable learning curve. I have left the placement feeling much more technically proficient and professionally competent than when I finished the second year of my course.

Some of the first things that I learned when I started the placement was how to read other people's code and how to mentally navigate a large code base. While these things may seem trivial to the established developer I found that they were essential skills that could only be gained by exposure to a professional development team.

When I started the placement I had only done a single module of C++ and I was self-taught to the beginner level in Python. After a year of being heavily exposed to both languages, I feel confident enough to call myself an intermediate in both languages. Besides gaining a better understanding of new languages through constant exposure, I have also gained experience with a variety of supporting libraries and tools. For example, I had not worked with the boost library before my placement which was used extensively throughout the project to provide C++11 features in a cross-platform way. I have become much more adept at using development tools such as git, gdb, cmake, and valgrind which I had little or no experience with before my placement.

Beyond enhancing my technical knowledge of my software development the placement has also provided me with great opportunities to better my "soft skills". Over the course of the 14 months I spent at RAL, we released a new version of Mantid 4 times. At three of the releases I was asked to present what I had been working on to the ISIS faculty which mainly consisted of instrument scientists. The content of what I presented varied, but I usually presented any changes to the indirect inelastic section, but I also presented more general topics on occasion (such as the algorithm history changes see section Improvements to Algorithm History Recording). I was also asked to present directly to the MSG during one of their monthly meetings and present any proposed core changes to Mantid during their biweekly development meetings. As I would not describe myself as a natural speaker, opportunities to present material to a variety of different audiences (both in size and demographic) proved to be immensely beneficial experience.

Talking to and gathering requirements from users was another skill that felt I developed well over the course of the year. I met with specific members of the MSG on a regular basis to pin down the requirements of the group. I also encouraged team members during the monthly meeting to contact me if they had any features or bug fix requests, many of whom did. This allowed me to quickly build a solid relationship with the scientists in the group who would regularly contact me with their issues. Being an informal working environment I ended up usually just going to the relevant persons office with a notepad and chatting about any issues/feature requests, before going anyway and implementing the request and getting it into the nightly build for them to play with. I felt that generally this system worked quite well. Users in many cases could physically show the issue they were having on their computer and it was faster than emailing back and forth. I felt this approach of regular contact also broke down the “them and us” wall between scientists and the development team.

Despite these many positive aspects, there were several low points to the placement that I felt could of been improved upon. While the development team I worked with was for the most part friendly and helpful I found that there appeared to be a lack of cohesion within the group. Team members seemed to work fairly independently of one another in their own areas which caused diversion of goals and a lack of communication which manifested itself as defects in the project. Most of the GUIs for individual experimental techniques are completely different from one another (not just in function, but in look and feel and interface object conventions). I felt that the development team could of benefitted from some basic team building activities or even the occasional social activity to try and build teamwork. I believe that the Mantid team was the only group at RAL who didn’t have coffee breaks together!

Leading on from this point I noticed over the course of the year that there appeared to be project divergence between the American development team the UK development team. It is worth noting that the instruments based at the SNS primarily use event data (where each individual neutron detection event is recorded individually) while ISIS mostly uses histogram data, which means that there fundamentally must be some difference between the two. However, as with individual techniques, there appears to be a completely different line of development undertaken for the Americans, with their own separate GUIs and routines, much of which could potentially be combined. Again, I believe this was down to a communication problem. The team only had a single biweekly development meeting via Bluejeans (which often got cancelled) and an ongoing Skype IM conversation as the means to contact one another. This meant that the people on both sides of the atlantic would often not communicate for days or weeks at a time.

I feel the simplest solution to this issue would be to increase the number of opportunities that members on both sides of the atlantic have to talk to one another. I would suggested upping the number of code reviews (at least one a week) and having a daily standup meeting with both sides in attendance over Bluejeans. I think this would help to keep both development teams better informed of the general direction of the project.

I also found there to be several issues in the specific area I was working in. The quality of the code base underlying the indirect geometry section of Mantid was generally poor. There was only minimal coverage of the code base with system tests and zero unit tests. A lot of the underlying python scripts and Fortran routines had been written by scientists and were not of production quality and no attempt had been made to refactor the code before integrating it into the Mantid framework. Docstrings, commenting and documentation was almost non-existent.

This proved to be incredibly difficult to work with. Large portions of my time were spent re-writing code to make it more reusable, maintainable, and better documented. When receiving new scripts from scientists I would aggressively refactor it and pester the writer for documentation of why the routine was useful and what it does (particularly if they could produce a paper outlining the method). I also began to write unit tests and system tests for any new routines. While the code base was still no where near the level of quality I was hoping for by the time I left, I did feel that these measures were slowly starting to make an impact on the quality of the code.

Finally, one further negative issue I had to deal with on several occasions was a lack of direction and conflict of interest between developers and instrument scientists. Several of the scientists within my

group approached me to request features that they would like to see or to suggest improvements. This is good because it means we write what the use wants, but can become burdensome when there is a lack of prioritisation of goals. I often found myself “spinning many plates” for many different task masters. Occasionally this led to a conflict either between what individual scientists want or between what scientists want and the general goals of the project. Towards the end of the year I felt we began to rectify this by having a weekly meeting where myself, my line manager (an MSG scientist) and a retired scientist who’s proffesion was QENs data analysis along with guests from both the development team and the scientific group. This allowed us to better outline what our goals should both in the short an long term, as well as increasing communication between all parties.

In summary, I feel that I had an excellent placement a ISIS. I was given responsibilities above and beyond anything I had imagined before starting my industrial year. I met and worked with some very interesting people and an unusual company. The learning curve was the steepest I’ve every experienced, but I enjoyed every moment. Despite several negative aspects to the project and making many mistakes along the way, I feel that I was able to learn from them and was able to pick myself up and carry on.

References

- [1] ISIS Neutron Source, “How ISIS works.” <http://www.isis.stfc.ac.uk/about/how-isis-works6313.html>, May-2014.
- [2] M. Project, “Mantid: Manipulation and analysis toolkit for instrument data,” 2013 [Online]. Available: <http://dx.doi.org/10.5286/SOFTWARE/MANTID>
- [3] D. S. Sivia, *Elementary scattering theory for x-ray and neutron users*. OUP Oxford, 2011.
- [4] F. Fernandez-Alonso and D. L. Price, *Neutron scattering fundermentals*, 1st ed. Academic Press, 2013.
- [5] V. Garcia-Sakai, M. Adams, W. Howells, and M. Telling, “The IRIS user guide,” Rutherford Appleton Laboratory UK, RAL-TR-2011-004, 2011.
- [6] M. Telling and K. Andersen, “The OSIRIS user guide,” Rutherford Appleton Laboratory UK, RAL-TR-2003-016, 2003.
- [7] J. Mayers and G. Reiter, “The VESUVIO electron volt neutron spectrometer,” *Measurement Science and Technology*, vol. 23, no. 4, p. 045902, 2012.
- [8] R. Pynn, “Neutron scattering: A non-destructive microscope for seeing inside matter,” in *Neutron scattering applications and techniques*, Springer, 2008.
- [9] ISIS Neutron Source, “Molecular spectroscopy group.” <http://www.isis.stfc.ac.uk/groups/molecular-spectroscopy/molecular-spectroscopy-3384.html>, May-2014.
- [10] Mantid Project, “Git workflow.” http://www.mantidproject.org/Git_Workflow, August-2014.
- [11] Mantid Project, “Workspace.” <http://www.mantidproject.org/Workspace>, August-2014.
- [12] Mantid Project, “Algorithms.” <http://www.mantidproject.org/Algorithm>, August-2014.
- [13] D. Sivia, C. Carlile, W. Howells, and S. Knig, “Bayesian analysis of quasielastic neutron scattering data,” *Physica B: Condensed Matter*, vol. 182, no. 4, pp. 341–348, 1992.
- [14] Open Genie, “Open genie homepage.” http://www.opengenie.org/Main_Page, August-2014.
- [15] J. Mayers and M. Adams, “Calibration of an electron volt neutron spectrometer,” *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, vol. 625, no. 1, pp. 47–56, 2011.

[16] Github, “Mantid project: Nested history detialed design document.” <https://github.com/mantidproject/documents/blob/master/Design/Nested%20History%20Detailed%20Design%20Document.docx>, August-2014.