

Industrial Year Report

Software Developer at ISIS Neutron Source

Samuel Jackson
University Of Aberystwyth
slj11@aber.ac.uk

October 2, 2014

Contents

1	Introduction	3
1.1	Introduction to Neutron Scattering	3
2	Organisational Environment	5
2.1	Organisation of STFC	5
2.2	Organisation of ISIS Neutron Source	5
2.3	Organisation of the Mantid Project Development Team	5
3	Technical and Application Environments	7
4	Description of Job Role and Work Done	9
4.1	Initial Training and Development Work	9
4.2	Indirect Bayes Interface	9
4.3	VESUVIO Calibration	10
4.4	Density of States Algorithm	11
4.5	Improvements to Algorithm History Recording	11
4.6	Development Reports	12
5	Critical Evaluation of Placement	14
	Bibilography	17

1 Introduction

This report details the industrial placement year I undertook as part of the Software Engineering MEng course at Aberystwyth university. The position of the placement was a junior role as part of the Mantid data analysis toolkit development team based at the ISIS facility at Rutherford Appleton Laboratory in Harwell, Oxfordshire, which is owned by the Science and Technologies Facilities Council (STFC). I held this post for a total duration of 14 months, with 12 months on the original contract and a 2 month extension.

ISIS is a world leading neutron and muon scattering facility. The facility operates two target stations and a 800 MeV pulsed proton synchrotron which acts as a source of neutrons for both neutron and muon time-of-flight spectroscopy [1]. Neutron and muon spectroscopy is used to probe the structure and dynamics of materials at the atomic level.

The Mantid project [2] is a free, open source application which aims to provide a single, unified application for the analysis of neutron and muon scattering data from facilities such as ISIS. The project is primarily developed by two teams of developers, one based at ISIS and one at the Spallation Neutron Source (SNS) at Oakridge laboratory in Tennessee, USA.

1.1 Introduction to Neutron Scattering

In order to fully understand the operation of ISIS and what the Mantid application is used to analyse, a small amount of background knowledge of the techniques used in neutron scattering experiments is useful along with some definitions of key terms used throughout this document.

Instruments at the ISIS facility are designed to probe the structure and dynamics of materials at the atomic scale. ISIS operates by accelerating protons to 84% the speed of light using a synchrotron. The resulting particles are then directed towards a tungsten target in one of the two target stations and from there towards samples in the spectrometers at the end beamlines connecting to the target station. The resulting collisions causes a scattering pattern which is detected by the spectrometers and analysed using programs such as Mantid.

All instruments at the ISIS accelerator operate using the time-of-flight technique. This is where the time for a neutron to travel from the source to an instrument's detector is accurately known. This value when combined with known parameters of the instrument (i.e. the length of the incident and final flight paths, the scattering angle θ , and the azimuth angle ϕ) can be used to determine interesting properties about the sample. The raw value measured by experiment is the time it takes a neutron to reach a detector in microseconds. In Mantid, this data is typically stored as a histogram with a count of the number of neutrons detected along the y-axis and the time-of-flight in microseconds on the x-axis.




Figure 1.1: A diagram of an indirect geometry spectrometer. M is the moderator, S is the sample, and D is the detector. L_1 is the length of the incident neutron flight path, L_2 is the final flight path, θ is the scattering angle, and ϕ is the azimuth angle. In indirect geometry spectrometers the final energy is fixed and is selected using a crystal analyser (typically pyrolytic graphite).

Broadly speaking, neutron scattering can be split into two categories: elastic and inelastic. Elastic scattering is where the final energy of a scattered neutron is equal to the energy of the incident neutron,

i.e. there is no transfer of energy to or from the sample. Inelastic scattering, which is the technique generally used by instruments belonging to the molecular spectroscopy group (MSG) that I was attached to (see section 2.3), is the more complex case where the energy of the incident neutron and the scattered neutron are not equal, i.e. there is a transfer of energy to or from the sample. From this transfer in energy and from known parameters of the instrument an instrument independent scattering function can be defined which provides a full model of the sample. This function is usually denoted as $S(Q, \omega)$, where Q is the momentum transfer and ω is energy transfer [3, 4].




Figure 1.2: Neutron scattering techniques by energy range.

Two important types of inelastic scattering are quasi-elastic neutron scattering (QENS) and deep inelastic neutron scattering (a.k.a neutron Compton scattering). Quasi-elastic neutron scattering is the case where the energy transfer is very close to zero and the scattering is almost elastic. This is typically performed on low energy spectrometers such as IRIS and OSIRIS [5, 6]. Deep inelastic neutron scattering (DINS) is the opposite case which uses extremely high energies ($>1\text{eV}$) and is used to measure the momentum distribution of atoms. This technique is still in its developmental phase, with the VESUVIO spectrometer at ISIS currently being the only instrument in the world capable of DINS [7].

Each of the spectrometers used by the MSG are of what is known as indirect geometry. This is where the final scattered energy is restricted to a particular wavelength (usually with a analyser crystal which will absorb all other wavelengths) and the incident neutron energy is varied. The value of the incident energy can then be found though the laws of the conservation of energy [3].

My work in with the Mantid team was centred round improving and maintaining the code base for analysis of data from indirect geometry instruments at ISIS. Mantid includes a framework for loading, processing and visualising data from the raw time-of-flight measurements. For a more in-depth introduction to neutron scattering the reader is directed to the references, in particular refs. [3, 4, 8] provide the best introductions to those unfamiliar with the subject.

2 Organisational Environment

2.1 Organisation of STFC

The Science and Technologies Facilities Council (STFC) is a UK based government funded body that carries out a wide variety of scientific research across a multitude of disciplines including particle physics, nuclear physics, space science and engineering, medical and biological sciences, and computational science. It was formed in 2007 by the merger of the Particle Physics and Astronomy Research Council (PPARC), the Council for the Central Laboratory of the Research Councils (CCLRC), and the Engineering and Physical Sciences Research Council (EPSRC) [9].

While the organisation is funded by the government, it is classified as a non-governmental body which acts as an umbrella organisation for an array of facilities based across the UK. These include (but are not limited to) the Central Laser Facility, Diamond Light Source (which is a publicly limited company of which STFC holds an 86% share), ISIS neutron source, and RAL Space; all of which are located at Rutherford Appleton Laboratory in Oxfordshire. The organisation also owns the Daresbury Laboratory located in Cheshire and the Chilbolton Observatory based in Hampshire.

The STFC's head office is located at Polaris House, Swindon, Wiltshire and is headed by the chief executive John Womersley. The purpose of STFC is the general organisation and management of the facilities under its control. In particular it is responsible for allocating budgetary and staffing allowances and liaising between government departments, particularly the department for business, innovation and skills which is the primary government body controlling STFC.

2.2 Organisation of ISIS Neutron Source

ISIS neutron source is a facility that is owned, operated and funded by the STFC. The organisational hierarchy of ISIS is headed by the director Prof. Robert McGreevy and includes several division heads for individual functional areas within ISIS such as the diffraction, spectroscopy and support, experimental operations, instrumentation, design, and accelerator divisions.

ISIS is also split into a number of research, operations, and experimental support groups. The computing group, of which the Mantid project is a part of, falls into the experimental support group. The computing group at ISIS is headed by Kevin Knowles. Other members of the computing group include facilities IT, business applications development, the ICAT data catalogue, and experimental controls.

While this description provides an overview of how the staffing of ISIS can be divided, in practice there tends to be a lot of cross over between sections of the organisation depending on an employee's skills and responsibilities. For example, while I was employed as part of the computing group, my line manager and senior manager were both members of the MSG whose interests I was responsible for within Mantid. However, I still reported to the Mantid project manager on a daily basis.

2.3 Organisation of the Mantid Project Development Team

The Mantid development team in the UK is a subset of the computing group. The Mantid group is headed by a single project manager (Nick Draper) who is based at ISIS where the project first started and is responsible for the overall management and direction of the project. The project is primarily split into two teams, one based at ISIS and one at the SNS in Oakridge, Tennessee. Both teams consist of a single lead developer and several senior developers who oversee the major technical developments and help to guide and manage the rest of the development team. The US team also has its own project manager but the project manager based at ISIS is in overall control.

Within Mantid, the project manager and the majority of the senior developers are actually contractors from Tessella Ltd. based in Abingdon, Oxfordshire. The rest of the development team are directly employed by ISIS, or in the case of the Americans, by the SNS. In addition there are also several

collaborators from facilities such as the Institut Laue-Langevin (ILL) and the Paul Scherrer Institute (PSI) who work remotely from their respective facilities.

While I was primarily situated within the development team based at ISIS, many developers within the team are generally also attached to a specific scientific group at the facility. In the author's case this was the ISIS molecular spectroscopy group (MSG) [10]. The members of the MSG (and other research groups) are effectively the users of Mantid and are the people who we gather requirements from, but this is typically an informal relationship and I worked closely with its members throughout the year.

3 Technical and Application Environments

The development team based in the UK has a single office located in the main office building of the ISIS facility. This office consists of approximately 13 workstations. The number of machines varied throughout the year depending on the level of staffing available to the project. Each of the machines were reasonably powerful 64-bit Dell workstations with between 8-16Gb of RAM and 8-16 core Intel i7 processors. Typical hard drive space for the machines was between 512Gb to 1Tb with the majority still being disk drives, but some of the newer machines had flash based storage.

The operating system that each machine ran was completely left to the preference of the developer, but it was recommended that the developer run one of the operating systems officially supported by Mantid for obvious reasons. In practice this meant that there was a good variety of developers using different platforms. I personally chose to run Ubuntu 12.13 as my operating system of choice for the majority of development work, with a dual partition running Windows 7 which could be swapped to when circumstances required. Other operating systems used by developers in the team included Windows 7 & 8, Mac OS X Mountain Lion and Mavericks, Red Hat Enterprise Linux 6, and Fedora 20.

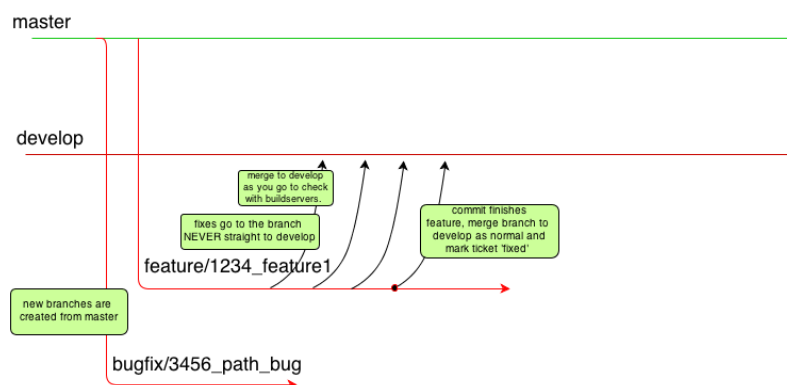


Figure 3.1: Diagram showing how the team used git branches to work on new features and bugfixes. Work is carried out on a branch from master, then merged to the develop branch for integration testing. Once integration testing is successful the ticket is marked as fixed. Source: [11]




Figure 3.2: Diagram showing how the team used git branches to test features/bugfixes. The ticket is merged with a local copy of master and tested by the tester. If testing fails the ticket is reopened. If testing passes the branch is merged with the master on GitHub and the feature branch deleted. Source: [11]

Apart from the workstations, the development team also had a collection of Jenkins build servers in order to support a continuous integration and testing workflow in conjunction with the Gitflow workflow [11]. See diagrams 3.1 and 3.2 for a description of how this workflow works. The build servers are jointly located at both ISIS and the SNS. At the start of the placement, the build servers for ISIS and the SNS were completely separate and located at different web address. Each individual build was run as a single job on the Jenkins build servers. Half way through my placement this was changed so that the servers were located at the same web address and the organisation of the build servers was changed to make use of matrix builds. This is where multiple builds are kicked off at the same time under a single umbrella job. For example the development branch matrix build would build the project and run the unit tests on each officially supported OS every time a new commit was pushed for integration testing.

Like the choice of operating system, the development software used by the team was flexible and open to developer preference. The project is built using the CMake build system on all supported platforms. On windows platforms the only supported compiler is Visual Studio 2012 or 2014 and most developers either chose to use the Visual Studio IDE, the Qt Creator IDE, or the Eclipse IDE. On Mac the Intel C++ compiler is used and typical IDEs are XCode, Eclipse, or Qt Creator. On Linux distributions the GNU compiler is the main supported compiler, with either Eclipse or Qt Creator used as the IDE for development. Many Linux developers are also happy to just use the make command to build the project from the command line. This approach is often used in conjunction with lightweight editors such as vim or sublime text. For interface development, Qt Designer was used on all platforms for anything more than the most trivial of jobs.

Unit tests are optionally built along side the project using a separate build target generated by CMake using the CxxTest unit testing framework. System tests are written in Python and make use of a collection of custom scripts loosely based on the unittest python module which makes use of the Mantid Python API. Debugging software used typically makes use of Visual Studio debugger on Windows, XCode/GDB on Mac and GDB on Linux distributions.

The Mantid application makes use of data files produced directly from neutron spectrometers. These files are collected on servers based at ISIS directly from the instruments themselves and are maintained by the scientific computing department. These provide both the instrument scientists and visiting scientists with direct access to their data. The development team also has access to these servers through network drives. On Windows operating systems access is provided through the in-built network drive capabilities. On Mac and Linux access is obtained through using Samba in conjunction with the SMB protocol. Copies of actual instrument data are frequently used as part of test scripts, especially in the case where the data required for the test cannot be easily simulated programmatically.

With regards to project management, the development team makes use of the git version control system and keeps all of the source code openly available to the public via Github. The Trac ticketing program is used to keep a record of the current development progress. The Trac set-up used also includes plug ins to automatically capture commits to branches on Github and update the relevant ticket with the commit information.

The Mantid documentation pipeline was overhauled during my time on placement. Originally the project used MediaWiki with the documentation for each class residing in the file that contains the code. This was changed to use Sphinx and accompanying reStructured text files for every part of our public python API. For the C++ code we continued to use regular Doxygen documentation. The website continued to use MediaWiki with a custom theme. As part of the upgrade examples of how to use the public API were added to the documentation. These are automatically run by Jenkins and their output is checked against reference results to help prevent the usage examples becoming out of date or broken. They also act as ersatz unit tests.

4 Description of Job Role and Work Done

As mentioned in section 2.3, the majority of my year was spent attached to the Molecular Spectroscopy Group (MSG) at ISIS. My role in the development team was to satisfy the computational data analysis requirements of the MSG within Mantid. This included involvement in every part of the development cycle, from gathering requirements from the users (the instrument scientists in the MSG) through to implementation of requested features, to testing and maintenance/bug fixing.

Before going into discussion about the work I carried out as part of my placement, it is useful to define a couple of core concepts used within the Mantid application: algorithms and workspaces. In Mantid a workspace [12] is an entity that contains a dataset. There are several types of workspace, but the most common is a MatrixWorkspace which is a collection of n X, Y, and E spectra either as a histogram or as point data. Each spectrum typically corresponds to a single detector but this is not necessarily always the case. An algorithm [13] is a class defined using the the Mantid algorithm framework that manipulates a workspace in some way (such as loading data into a workspace or rebinning a workspace). A common analogy is that workspaces are the ‘nouns’ of Mantid while algorithms are the ‘verbs’.

4.1 Initial Training and Development Work

My first couple of weeks at ISIS were spent running through some training exercises setup by the development team. These were relatively simple programming problems designed to gauge my existing understanding of C++ and Python. After spending a week completing these introductory activities, I had an informal code review with the lead developer who examined the good and bad points of what I had written before giving me an introduction to the Mantid application and workflow of the team. The following week was mostly spent setting up my workstation for development before starting my first ticket from the Trac ticket system. My first ticket involved writing a couple of algorithms which would take some histogram data and perform smoothing and interpolation of the workspace using the cubic spline routines in the GSL library, a feature that had been requested by one of the MSG scientists. From this point onwards I spent the majority of my time working on tickets for the MSG.

4.2 Indirect Bayes Interface

Apart from maintenance and minor feature requests my first real major piece of development work was to write a GUI for a collection of Bayesian fitting routines that had been provided by one of the scientists in the group. These routines are used as part of quasi-elastic neutron scattering analysis to determine the shape of the scattering function present in the dataset by using Bayesian methods to compute the likelihood of a given model. These methods are based on the original routines described by Sivia [14].

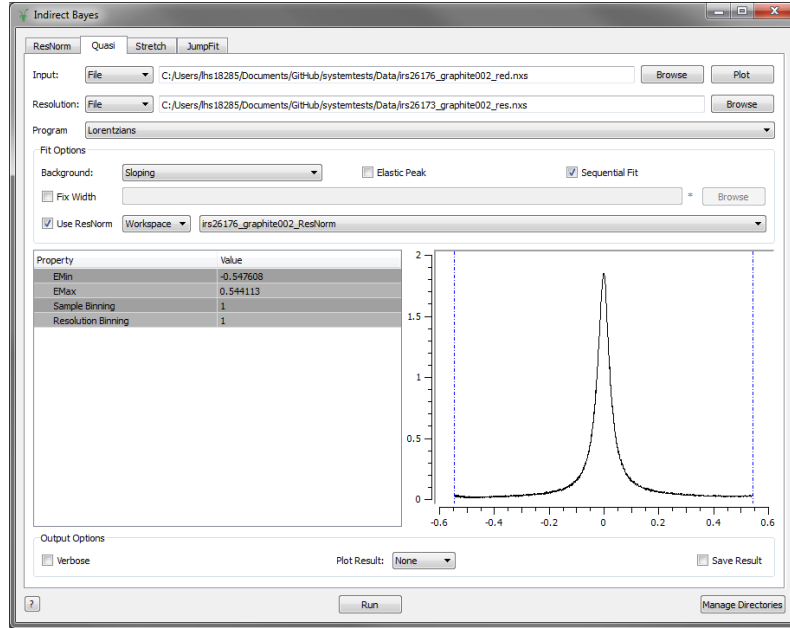


Figure 4.1: Screen capture of the Indirect Bayes QENs analysis interface in Mantid.

The purpose of the interface was to provide an easier way for users to set up and run a fit to their data than by using the existing scripts which required many parameters to be specified. This was fiddly and often led to the user inputting incorrect combinations of parameters causing the script to throw an error or even crash Mantid.

The solution to this was to add another custom interface under the indirect section of the application (there were already three others in place). This consisted of a collection of C++ classes, created using the Qt framework, one for the parent window of the GUI and one for each of the individual routines on the interface which were implemented as separate tabs on the interface. Each of the routines were refactored from the original code base and required updated system tests, which were implemented at the same time as the new interface.

4.3 VESUVIO Calibration

Another major piece of work I was involved with during my placement was porting some calibration scripts for the VESUVIO instrument [7] which were originally written as Fortran programs for the older OpenGenie application [15]. As part of this project, I created a new implementation for Mantid based on the original calibration procedures described in ref. [16].

This required rigorous testing by the scientist responsible for the instrument and was built slowly over several of months. The Mantid implementation was designed to be radically different from the original. The new version was built using the Mantid concepts of algorithms and workspaces and was intended to be more maintainable than the original. I also wrote a couple of unit test suites used to check the quality of the calibration using several different sample materials at the same time.

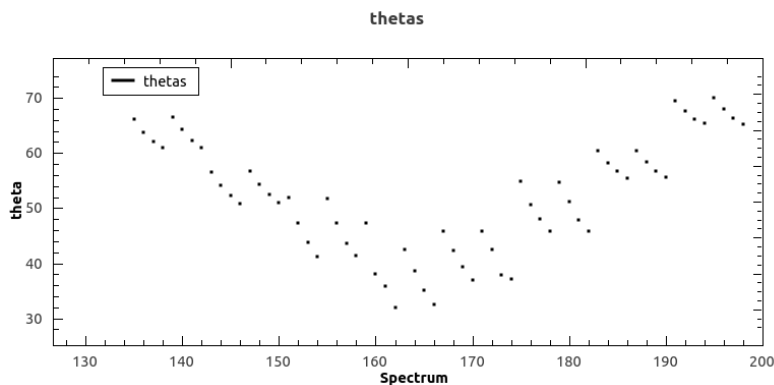


Figure 4.2: *Plot of the calibrated values for θ from a Pb sample output from the VESUVIO calibration program. Compare this with the θ plots in reference [16].*

The end product was a pair of Mantid algorithms, one for fitting the appropriate line shape to each peak in every spectrum of a sample run specifically measured for calibration with the Mantid python API and one which calculated the calibration parameters for the instrument using the parameters of the fit obtained from the first algorithm.

4.4 Density of States Algorithm

During my time with the MSG, it became apparent that one of the major areas for development that had not been explored was the implementation of support for the simulation of a neutron scattering experiment. This was an area which my supervisor within the MSG (who has a background in molecular simulation) was keen to expand because comparison with simulation is one of the most common and useful techniques for analysing experimental data.

One of her requests was to add support in Mantid for loading the results of a simulation produced using the CASTEP code [17]. My supervisor already had a perl script which could do the calculations she required through the command line, but it was infinitely more convenient to have this functionality within Mantid.

This involved the creation of a new algorithm that could read the files output from CASTEP, which contained a list of frequencies predicted by theory for a sample. With this information, the density of states could be calculated and a workspace containing the simulated spectrum could be created. Functionality was also added to calculate the infrared and Raman spectra from data in the files. Later in the placement she requested additional functionality be added to the algorithm to calculate the partial density of states from CASTEP code.

4.5 Improvements to Algorithm History Recording

Occasionally I became involved in the more general development of the Mantid framework. The largest piece of general development I was asked to do was to improve the algorithm history system. In Mantid, when algorithms get executed on a workspace a record of the algorithm run and the parameters used to execute the algorithm are stored on the workspace object. In theory this can provide a detailed record of everything that has happened to the workspace from the moment the data entered the application. In practice this feature had only been partially implemented and had slowly become more redundant/broken as development moved on.




Figure 4.3: Screen capture of the algorithm history GUI showing nested algorithms used on a workspace from the *INTER* spectrometer.

My task was to overhaul this feature to not only properly provide the features mentioned above, but also to add functionality to capture nested algorithm history, that is, the history of an algorithm operating on a workspace but also the history of any ‘child’ algorithms run as part of that algorithm. There was also a requirement to make the workspace completely reproducible from its history alone and that the full nested history could be examined using a GUI from within Mantid.

As this consisted mostly of core system changes a much more rigorous design phase was required. A draft of a design document already existed for the algorithm history which I presented to the development team (both in the UK and the states) and was roundly rejected for being too vague. The next step was to prototype out some of the more unclear features such as how history should be saved to file which needed to be a compromise between space and time efficiency, how to represent nested history at runtime (e.g. `boost::graph`, `std::list` of children, `std::set` of children?), and how nested history records should be stored on disk.

After prototyping the system, I created a new design document with the findings from the prototypes and my suggested proposals for change [18]. This document was reviewed by the senior members of the development team and with a few minor changes was approved. I then spent the next month or so implementing the changes outlined in the design document. Along the way we encountered several issues that were not covered by the design document, mostly linked technical issues caused by the fact that workspaces can be created without a string name and that providing a temporary name to them on the fly can cause many issues with reproducing an executable script.

Despite these setbacks, I managed to get the majority of the core changes into Mantid in time for the next release. However there were still several parts of the system left to the next release cycle, such as compressing the size of large algorithm properties (such as when arrays are passed as parameters to an algorithm), which was left to my successor.

4.6 Development Reports

One final piece of work I was asked to complete during my time at ISIS was to produce two major reports. The first [19] was to provide a snapshot of the current progress of development of the indirect framework. This report gives an overview of the entire framework as of release 3.1 including a description of the theory behind the routines and detailed description of how each one works. Suggested ideas for future development are also listed.

The second report [20] was a basic user manual for the routines still under development for the VESU-

VIO spectrometer and provides a description of how to analyse data from the instrument with the implementation as it stands. Both of these reports are publicly available for free via the RAL ePubs system.

5 Critical Evaluation of Placement

My overall experience with the placement was a positive one. ISIS provided me with a level of responsibility far exceeding anything I expected. For the most part, I felt that I was a valued member of the development team and that I was treated as an equal rather than as an intern. I feel that my practical understanding of software development has increased tenfold, particularly over the first few months which I found to be a steep but enjoyable learning curve. I left the placement feeling much more technically proficient and professionally competent than when I finished the second year of my course.

One of the first things that I learned when I started the placement was how to read other peoples code and how to mentally navigate a large code base. While these skills may seem trivial to the established developer I found that they only be gained by exposure to a large project with multiple developers.

When I started the placement I had only taken a single module in C++ and I was self taught to the beginner level in Python. After a year of being heavily exposed to both languages, I feel confident enough to call myself an intermediate in both languages. Besides gaining a better understanding of new languages though constant exposure, I have also gained experience with a variety of supporting libraries and tools. For example, I had not worked with the boost library before my placement which was used extensively throughout the project to provide C++11 like features in a cross platform and backwardly compatible way. I have also become much more adept at using development tools such as Git, GDB, CMake, and Valgrind which I had little or no experience with previously.

Beyond enhancing my technical knowledge the placement has also provided me with great opportunities to better my ‘soft skills’. Over the course of the 14 months I spent at RAL, we released a new version of Mantid 4 times. At three of the releases I was asked to present what I had been working on to the ISIS faculty which mainly consisted of instrument scientists with limited knowledge of the development process. The content of what I presented varied, but I usually presented any changes to the indirect inelastic section, but I also presented more general topics on occasion (such as the algorithm history changes, see section 4.5). I was also asked to present directly to the MSG during one of their monthly meetings and to present any core changes to Mantid during their biweekly development meetings. As I would not describe myself as a natural speaker, opportunities to present material to a variety of different audiences (both in size and demographic) proved to be an invaluable experience.

Talking to and gathering requirements from users was another skill that felt I developed well over the course of the year. I met with members of the MSG on a regular basis to pin down the requirements of the group. I also frequently encouraged MSG members to contact me if they had any features or bug fix requests, many of whom did. This allowed me to quickly build a solid relationship with the scientists in the group who would regularly contact me with feedback on their experiences with the application. Being an informal working environment, requirements gathering generally consisted of going to the relevant person’s office with a notepad and having a short chat about any issues/feature requests, before going anyway and implementing the requests and getting it into the nightly build for them to play with. I felt that generally this system worked quite well. Users in many cases could physically show me the issues they were having and it was faster than emailing back a forth. I felt this approach of regular contact also broke down the “them and us” wall between scientists and the development team.

Despite these many positive aspects there were several low points to the placement that I felt could of been improved upon. While the development team I worked with was for the most part friendly and helpful I found that there appeared to be a lack of cohesion within the group. Team members seemed to work fairly independently of one another in their own areas which seemed to cause a diversion of goals and a lack of communication which manifested itself as defects in the project. Most of the GUIs for individual experimental techniques are completely different from one another (not just in function, but in look and feel and GUI conventions). I felt that the development team could of benefitted from some basic team building activities or even the occasional social activity to try and build teamwork and get them talking to one another. I believe that the Mantid team was the only group at RAL who didn’t have coffee breaks together!

Leading on from this point I noticed over the course of the year that there appeared to be project divergence

between the American development team and the UK development team. It is worth noting that the instruments based at the SNS primarily use event data (where each individual neutron detection event is recorded individually) while ISIS mostly uses histogram data, which means that there fundamentally must be some difference between the two. However, as with individual techniques, there appears to be a completely different line of development undertaken for the Americans, with their own separate GUIs and routines, much of which could potentially be combined. Again, I believe this was down to a communication problem. The both teams only had a single biweekly development meeting via BlueJeans (which often got cancelled) and an ongoing Skype IM conversation as the means to contact one another. This meant that some people on both sides of the atlantic would often not communicate for days or weeks at a time.

I feel the simplest solution to this issue would be to increase the number of opportunities that members on both sides of the atlantic have to talk to one another. I would suggested upping the number of code reviews (at least one a week) and having a daily standup meeting with both sides in attendance over BlueJeans. I think this would help to keep both development teams better informed of the general direction of the project.

I also found there to be several issues in the specific area I was working in. The quality of the code base underlying the indirect geometry framework of Mantid was generally poor. There was only minimal coverage of the code base with system tests and zero unit tests. A lot of the underlying Python scripts and Fortran routines had been written by scientists and were not of production quality and no attempt had been made to refactor the code before integrating it into the Mantid framework. Docstrings, commenting and documentation was almost non-existent. I believe this was caused by years of placement students being left to look after the section on there own, some of whom were obviously less competent than others, resulting in disorganised and fractured code base.

This proved to be incredibly difficult to work with. Large portions of my time were spent re-writing code to make it more reusable, maintainable, and better documented. When receiving new scripts from scientists I would aggressively refactor it and pester the writer for documentation of why the routine was useful and what it does (particularly if they could produce a paper outlining the method). I also began to write unit tests and system tests for any new routines. While the code base was still no where near the level of quality I would expect by the time I left, I did feel that these measures were slowly starting to make an impact on the quality of the code. This taught me that unit tests give a developer confidence to make changes and how to refactor properly.

Finally, one further negative issue I had to deal with on several occasions was a lack of direction and the conflict of interest between developers and instrument scientists. Several of the scientists within my group approached me to request features that they would like to see or to suggest improvements. This is good because it means we write what the user wants, but can become burdensome when there is a lack of prioritisation. I often found myself ‘spinning many plates’ for many different task masters. Occasionally this led to a conflict either between what individual scientists want or between what scientists want and the general goals of the project.

Towards the end of the year I felt we began to rectify this by having a weekly meeting between myself, my line manager (an MSG scientist) and a retired scientist who’s profession was QENs data analysis along with guests from both the development team and the scientific group. This allowed us to better outline what our goals should both in the short an long term, as well as increasing communication between all parties.

All of these issues really outlined to me the importance of good communication within a development team. While these issues were not so horrendous as to prevent the team from delivering the software, I really felt that even the smallest amount of miscommunication can lead to bad engineering and long delays.

In summary, I feel that I had an excellent placement a ISIS. I was given responsibilities above and beyond anything I had imagined before starting my industrial year. I met and worked with some very interesting people and worked for unusual company. The learning curve was the steepest I’ve ever experienced, but I enjoyed every moment. Despite several negative aspects to the project and making many mistakes along

the way, I feel that I was able to learn from them all.

Bibilography

- [1] ISIS Neutron Source, “How ISIS works.” <http://www.isis.stfc.ac.uk/about/how-isis-works6313.html>, May-2014.
- [2] Mantid Project, “Mantid: Manipulation and analysis toolkit for instrument data,” 2013 [Online]. Available: <http://dx.doi.org/10.5286/SOFTWARE/MANTID>
- [3] D. S. Sivia, *Elementary scattering theory for x-ray and neutron users*. OUP Oxford, 2011.
- [4] F. Fernandez-Alonso and D. L. Price, *Neutron scattering fundementals*, 1st ed. Academic Press, 2013.
- [5] V. Garcia-Sakai, M. Adams, W. Howells, and M. Telling, “The IRIS user guide,” Rutherford Appleton Laboratory UK, RAL-TR-2011-004, 2011.
- [6] M. Telling and K. Andersen, “The OSIRIS user guide,” Rutherford Appleton Laboratory UK, RAL-TR-2003-016, 2003.
- [7] J. Mayers and G. Reiter, “The VESUVIO electron volt neutron spectrometer,” *Measurement Science and Technology*, vol. 23, no. 4, p. 045902, 2012.
- [8] R. Pynn, “Neutron scattering: A non-destructive microscope for seeing inside matter,” in *Neutron scattering applications and techniques*, Springer, 2008.
- [9] House of Commons Science and Technology Committee, “Office of science and innovation: Scrutiny report 2005 and 2006.” ISBN 0-215-03350-7, 2007.
- [10] ISIS Neutron Source, “Molecular spectroscopy group.” <http://www.isis.stfc.ac.uk/groups/molecular-spectroscopy/molecular-spectroscopy-3384.html>, May-2014.
- [11] Mantid Project, “Git workflow.” http://www.mantidproject.org/Git_Workflow, August-2014.
- [12] Mantid Project, “Workspace.” <http://www.mantidproject.org/Workspace>, August-2014.
- [13] Mantid Project, “Algorithms.” <http://www.mantidproject.org/Algorithm>, August-2014.
- [14] D. Sivia, C. Carlile, W. Howells, and S. Knig, “Bayesian analysis of quasielastic neutron scattering data,” *Physica B: Condensed Matter*, vol. 182, no. 4, pp. 341–348, 1992.
- [15] Open Genie, “Open genie homepage.” http://www.opengenie.org/Main_Page, August-2014.
- [16] J. Mayers and M. Adams, “Calibration of an electron volt neutron spectrometer,” *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, vol. 625, no. 1, pp. 47–56, 2011.
- [17] CASTEP, “CASTEP homepage.” <http://www.castep.org>, September-2014.
- [18] Github, “Mantid project: Nested history detailed design document.” <https://github.com/mantidproject/documents/blob/master/Design/Nested%20History%20Detailed%20Design%20Document.docx>,
- [19] S. Jackson, “An overview of the recent development of indirect inelastic data analysis in Mantid,” Rutherford Appleton Laboratory UK, RAL-TR-2014-010, 2014. September-2014.
- [20] S. Jackson, “VESUVIO Data Reduction and Analysis in Mantid,” Rutherford Appleton Laboratory UK, RAL-TR-2014-009, 2014. September-2014.