

CS26110 Assignment 1

Samuel Jackson - slj11@aber.ac.uk

November 13, 2012

1 Sliding Tile Puzzle

1.1 Node Representation

The problem of representing each node in the sliding tile puzzle within a search tree could be solved by using a fixed length array of size 7 where each element in the array represents a single tile in the sliding puzzle. The array could either store a list of characters of either B, R or space (representing the blue, red and empty slot respectively). This representation would be useful when outputting the resulting node to the user. Alternatively the problem could be represented using a trit array (with values 0, 1 and 2) where each trit represents a character in the array (or a space).

The reason why I believe an array to be a good choice for the representation of a node is that regardless of what symbol is used to signify a red, blue or space tile, the indices of the array give a handle to the tiles position in the grid. The indices of the array can then be used to swap the space with red and blue tiles. A tile's position in the array can also be used to calculate whether it is close enough to move into the empty space and calculate the cost of the move.

1.2 Initial State Tree

The diagram below shows all of the initial valid moves from the start state of the sliding puzzle game. The cost of the moves are shown alongside the line branching from the start state.




Figure 1: All possible moves from the start state, including the cost of the moves.

1.3 Uninformed Search Strategy

An uninformed search strategy could be used to solve the sliding puzzle providing that we have enough time and memory resources to represent and traverse the tree of valid moves. Although it is not specified by the rules of the game, it is obviously desirable that whatever search strategy we adopt to solve this puzzle should attempt to give us an optimal solution. It is also desirable that the search strategy is a complete one, as we know from the rules of the puzzle that there will always be more than one distinct solutions to the sliding puzzle because the end placement of the space does not matter.

I would suggest that the best uninformed search technique for this problem would be to use a uniform cost search to traverse the search space. There are several reasons why this is a good choice for the given problem. Firstly, uniform cost search takes into account the path cost of each edge between nodes in the search space. This means that the algorithm will try to take the lowest cost path between nodes in order to find a solution. Because of this it will generally prefer to take paths which do not involve a “double jump” to find the solution. Secondly, uniform cost search is guaranteed to find a solution which is both complete and optimal, so we are guaranteed to get the shortest, lowest cost path from the initial configuration to the end state.

Comparing this search strategy to other search techniques provides further support for my proposed use of uniform cost search. Both depth-first and depth-limited search techniques are not guaranteed to return an optimal solution. Also, depth-limited search is not guaranteed to return any solution at all, despite the fact that the search tree will always contain a number of different solutions. Clearly these options do not provide a viable solution to the given problem.

Other available search strategies include using a breadth-first search or a bi-directional search. Uniform cost search is effectively just a breadth-first search, except that uniform cost search considers path weights and breadth-first search does not. If all of the path weights were the same positive number then uniform cost search would exactly mimic a breadth-first search. Therefore on average we will get better performance out of a uniform cost search ($O(b^{1+C^*/\epsilon})$ time where C^* is the cost of the optimal solution), and in the worst case its performance will equal that of a breadth-first search ($O(b^d)$ time). Because of this, breadth-first search only guarantees to find us the shallowest solution, but not necessarily the lowest cost path to a solution.

A bi-directional search offers another option for consideration for an uninformed search strategy. Bi-directional search could be used to traverse both from the start configuration and at the same time traverse backwards from one of the end states. However, immediately there is a problem with this approach. Seeing as there are multiple ways to successfully complete the puzzle, how do we decide which end configuration to start from? One end configuration may yield a shorter path compared with another configuration and is therefore more desirable. Bi-directional search offers better time complexity compared with a breadth-first search as it effectively cuts the number of nodes that need to be expanded in half. However, it still does not consider the weight on any of the edges. Because of this, uniform cost search can potentially provide us with a better solution compared with a bi-directional search.

One final search strategy that could be considered is the use of the iterative deepening depth-first search. This will have a lower space complexity than uniform cost search, but still exhibits much of the same issues as a breadth-first search. Mainly that it does not make any attempt to consider the weight of the edge between two nodes, and therefore is not guaranteed to find the lowest cost solution, although it will find a complete one.

In conclusion, I believe that a uniform cost search is the clear uninformed search technique to be used to solve the sliding tile puzzle as it offers us a complete and optimal solution in a reasonable time and space complexity. However, traverse of the search space could be dramatically improved through the use of a heuristic as discussed in the next section.

1.4 Search Heuristic¹

One potential heuristic that could be used to solve this problem would be to take the current configuration of the sliding tile puzzle and assign to each tile a weight according to how many of the opposing coloured tile is currently on the wrong side of this tile. Each of these weights could then be summed and divided by two to give us a heuristic estimate of how far away the current configuration is from the target configuration. The reason that it must be divided by two is that each time two tiles swap, the weight on each tile will decrease

¹This section of the document answers questions 1(d) and 1(e) of the assignment together.

by one, leading total reduction of two. Summing the total cost of the tiles without dividing by two to account for this will lead to an overestimate and therefore render this heuristic inadmissible. This heuristic allows us to quantify how many tiles are out of order without caring about where specific tiles end up (as would be the case if we use a distance heuristic such as the Manhattan heuristic).

In this heuristic, all of the tiles initially have a weight of three assigned to them. This is because the start configuration has all of the blue tiles to the left of the red tiles and all of the red tiles to the right of the blue tiles. Hence, for each tile in the initial configuration, each has all three of the opposing colour on the “wrong” side and therefore its weight is three. As the tiles are moved around, some of them will jump over one another. At this point two of tiles are now closer to the solution than their counterparts and their weights will be modified accordingly as they will now only have two tiles on the wrong side of them. Weights of each tile would continue to be reduced as the algorithm progresses until eventually all tiles are in the correct configuration and the weights on each tile are zero.

The following diagram visually outlines part of a branch in the search space from the initial configuration where tiles might move and how their weights are adjusted accordingly. In this diagram, the subscripts of each of the tiles represent the current weight that tile has. Shown next to the node is the sum of the weights (our heuristic estimate) for that particular configuration.




Figure 2: An example of how the heuristic estimate changes as tiles are moved.

I believe this heuristic to be a good choice for this problem for several reasons. Initially, I considered that this puzzle presents certain characteristics that are closely related n-puzzle problems. Both problems involve sliding tiles around a fixed sized board and they both have only one space which can be moved into at any one time. Traditionally, one of the best heuristics used to solve a n-puzzle is the Manhattan distance heuristic, which I initially thought could be applied this problem. However there are several differences that mean it would have produced sub optimal results.

The key difference between this puzzle and a n-puzzle is that tiles in this puzzle can “jump” over two tiles at once, which incurs a different path cost compared to jumping over a single tile or moving into an open slot. Another key difference is that it does not matter what order the tiles end up in, providing that all the blue tiles end up on the opposite side of all the red tiles. This means that the Manhattan distance is not exactly the best heuristic for this problem, as I initially thought, because each tile does not have to end up in a specific

slot in the end configuration and therefore by using the Manhattan distance heuristic it would incur more processing than is actually required.

Subsequently while I believe that a Manhattan distance heuristic could be of value to a search algorithm used to solve this problem, I decided to scrap the idea in favor of a different technique that I have presented here. My suggested heuristic will given an estimate of the total distance left until we reach a solution, while not worrying about a tile moving to a specific slot in the puzzle.

I would like to suggest that the heuristic function outlined in this section is admissible, as it never overestimates the perfect solution to the problem. At best it only ever gives a heuristic estimate that is equal to the exact number of steps required to solve the problem. Otherwise the heuristic will underestimate the distance required to get from the initial configuration to a solution. This means that the heuristic function will always lead us towards a complete and optimal solution, rather than overestimating the path cost and returning a sub optimal solution. We could also combine this heuristic with an A* search to take account of the weights on the edges between states. This would then find an optimal solution through the lowest cost path to the solution while looking at a minimal number of states.

However, there are a few downsides to this heuristic. One downside is that it may underestimate the distance to the solution too much during some configurations of the puzzle. In this case, the heuristic offers little assistance to any search algorithm traversing the search space and will look at more states than if the estimate were more accurate. It may also be the case that certain moves do nothing to lower the heuristic estimate, such sliding a tile into the empty slot without jumping over anything. This would not lower the heuristic estimate and therefore the heuristic would not provide any additional information in this circumstance. However, despite these limitations, the heuristic would still offer some assistance in guiding a search algorithm towards an optimal goal state faster than an uninformed search technique.

2 Vehicle Grid Puzzle

2.1 Problem Definition

1. **Initial State** - The initial state of this problem starts with a set of n vehicles arranged in an $n \times n$ grid so that they line up along the “top” side of the grid, with each vehicle being in the first row of the grid and in the i -th column.
2. **Goal State/Test** - We have reached the goal state if all of the vehicles are in reverse order from there starting positions (i.e. each tile is now at column $n - i + 1$ where i is the starting index of the agent) on the “bottom” (n^{th}) row of the grid. We can test this by checking row n of each column to see if it contains vehicle $n - i + 1$.
3. **Actions** - An agent in this state space can move one square in any of the four directions defined in taxicab geometry providing that the square is not occupied by another agent or that the given direction takes the agent off of the grid. An agent may also chose not to move at all from its current position during a given turn.
4. **Path Cost** - The puzzle description does not outline that there is any specific cost associated with moving an agent to a new square relative to there direction, so it can be assumed that every move has an equal cost associated with it. Therefore each move an agent makes can be defined as having a cost of 1. As it is undesirable for each of the agents to not move, this will also have to incur a cost of 1. Therefore the total path cost of a solution can be fully quantified by how many total steps (in which each vehicle may or may not move) it takes to get all vehicles to reach the goal state.

2.2 Branching Factor

The approximate branching factor of the vehicle grid problem is 5^n . This is because for each state in the search space, there will be approximately five to the power of n possible valid choices leading to other valid states. This is because for every node in the search space, each vehicle can potentially “move” in five different directions: up, down, left, right and staying stationary. The reason why the branching factor is raised to n (the number of vehicles on the grid) is because for every node, each vehicle on the grid will make one of its

moves. Therefore, we can choose one move for each of the vehicles on the grid in any combination and so the product rule can be used to calculate the total number of combinations. In the case of a grid of size five where each vehicle can potentially move in up to five directions this is approximately $5 \times 5 \times 5 \times 5 \times 5 = 5^5 = 3125$ possible moves from each node.

2.3 Minimum Number of Steps

The minimum number of steps required for one vehicle to move from its current position at (r_i, c_i) , where i is the index of a column or row, can be expressed as the mathematical function:

$$h_i(n) = |r_i - n| + |c_i - (n - i + 1)|$$

Which gets us the vehicle i 's Manhattan distance from its current position in the grid to its finishing position at $n - i + 1$. This formula works by getting the absolute difference between the vehicles current row and the target row. It then sums this with the absolute difference between the current column and the target column for this vehicle $(n - i + 1)$. As an example, lets say that a vehicle of index 3 is on an empty 5×5 grid at position $(3, 4)$ then according to the formula above, the distance it would have to travel to reach its goal state would be a minimum of 3 moves. Note that this formula will work with any $n \times n$ grid; not just for a grid of size five.

2.4 Vehicle Grid Heuristic

The function outlined in the preceding section could be used as part of a heuristic function to help aid a search algorithm to solve the vehicle grid problem. The underlying principle of function h_i is that it calculates the Manhattan distance for the current vehicle to the goal state. In other words, based on the current position of the vehicle, it calculates how many steps, at minimum, need to be taken for the vehicle to reach its destination square. A heuristic function can use this to estimate the minimum number of moves required to get all of the vehicles from their current state to the goal state. To obtain a heuristic estimate, we simply apply the function outlined in section 2.3 to each of the vehicles in a given state and then sum the results as shown by the following formula:

$$h(n) = \sum_{i=1}^{i=n} h_i(n)$$

Where $h(n)$ is the heuristic cost estimate of moving all the vehicles from there current state on the grid to there goal states. Of course, this heuristic is not perfect, as it only takes into account how many moves need to be taken by a vehicle to get to its goal state on an empty grid. The Manhattan distance heuristic in no way accounts for other vehicles getting in the way of the current vehicle's path and therefore forcing it to either not move this step or to take a different (sub optimal) direction. However, this feature does mean that the heuristic is admissible as it will always be less than or equal to the actual total distance required for a vehicle to reach its goal state. Therefore the best search technique to use for this problem would be to use an A* search that takes account of the total distance traveled by each of the given vehicles for its path cost and uses the heuristic function $h(n)$ in order to find a optimal route through the search space.